

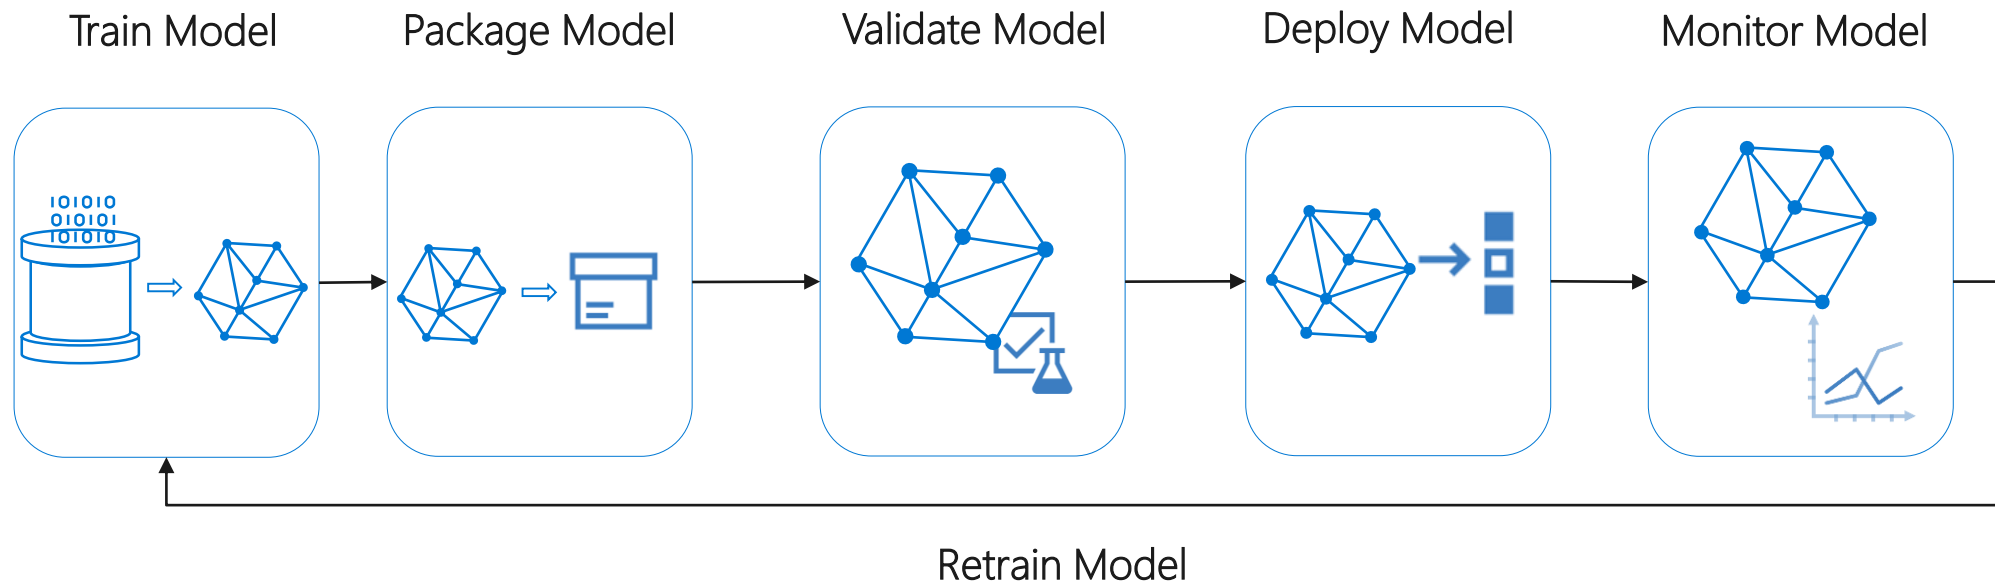


DevOps for ML MLOps



What does the Machine Learning Lifecycle look like?

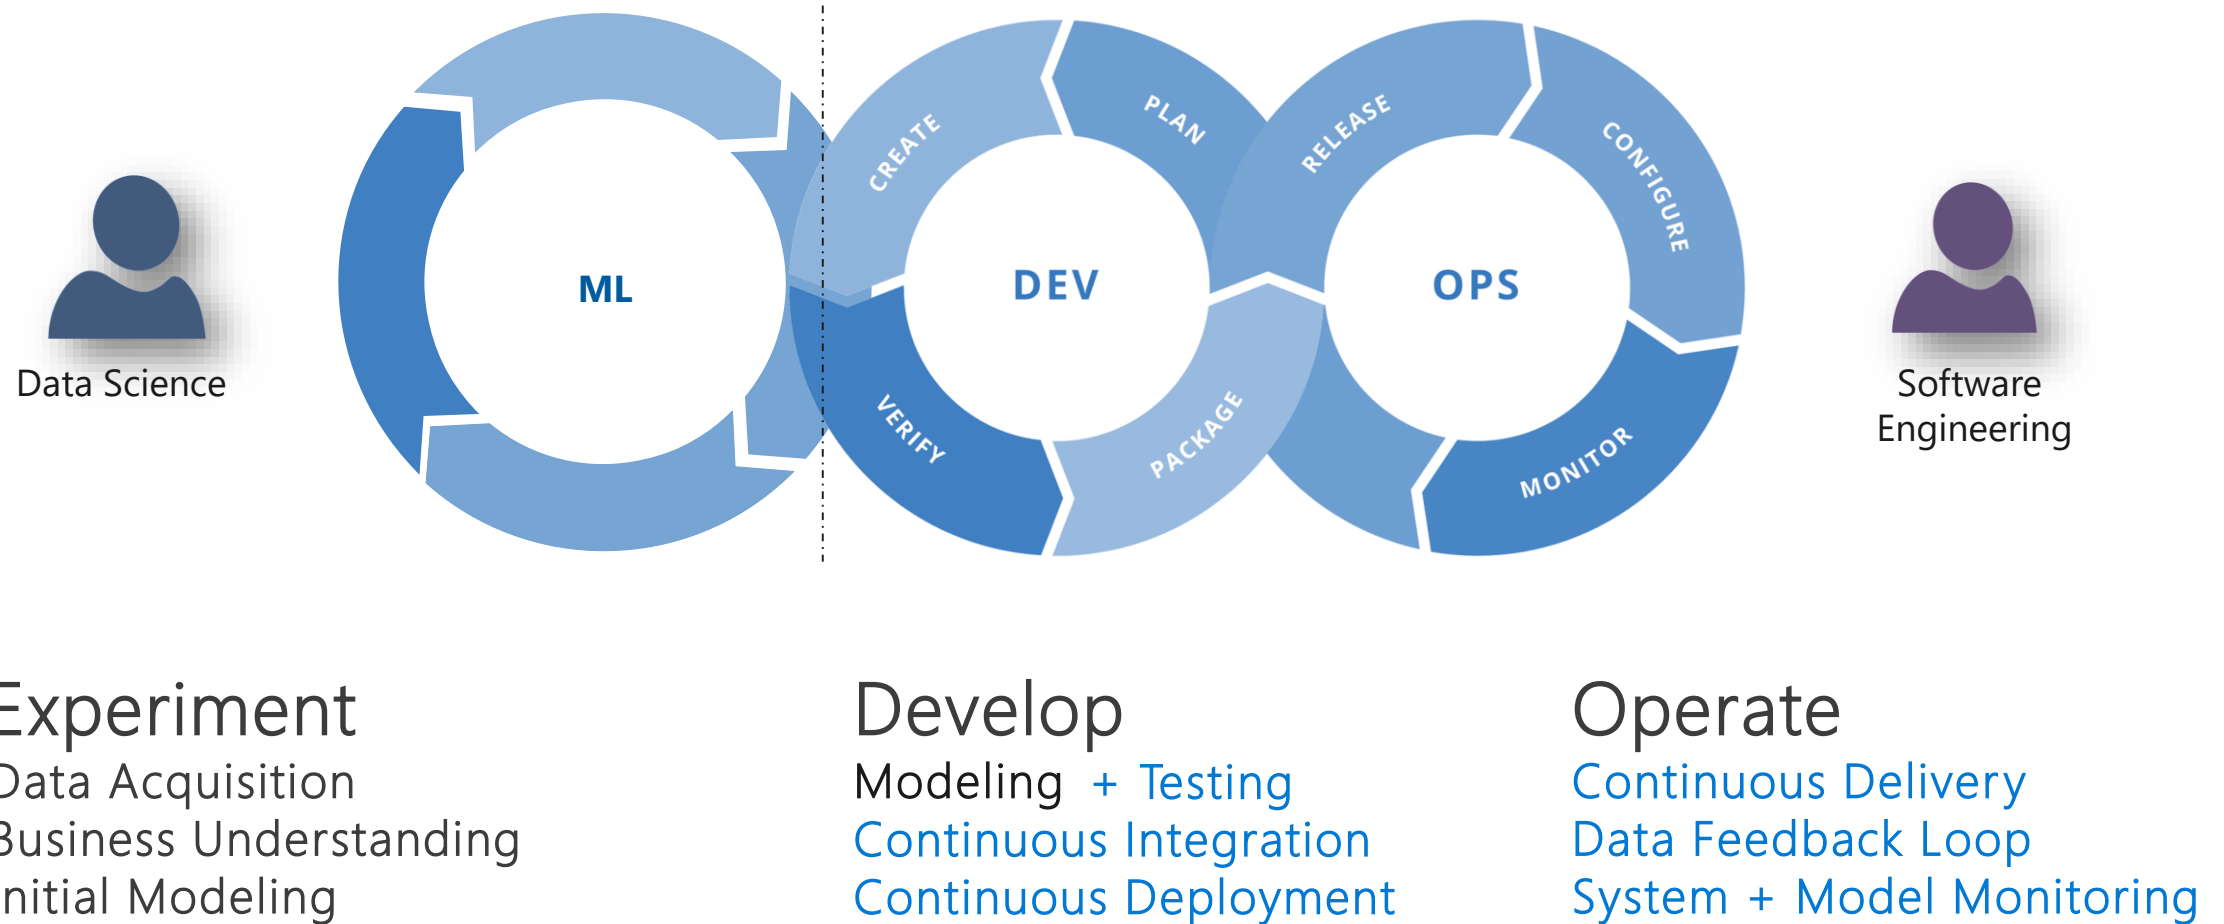
- **Develop & train model** that solves a real business problem
- **Package model** so you can use it somewhere else
- **Validate model behavior** – functionally, in terms of responsiveness, in terms of regulatory compliance
- **Deploy model** – use the model to make predictions
- **Monitor model** behavior & business value, know **when to replace / deprecate a stale model**



What is MLOps?

MLOps = ML + DEV + OPS

Help bring ML workflows to production




How is MLOps different from DevOps?



Data/model versioning != code versioning - how to version data sets as the schema and origin data change



Digital audit trail (lineage)
requirements change when dealing with code + data

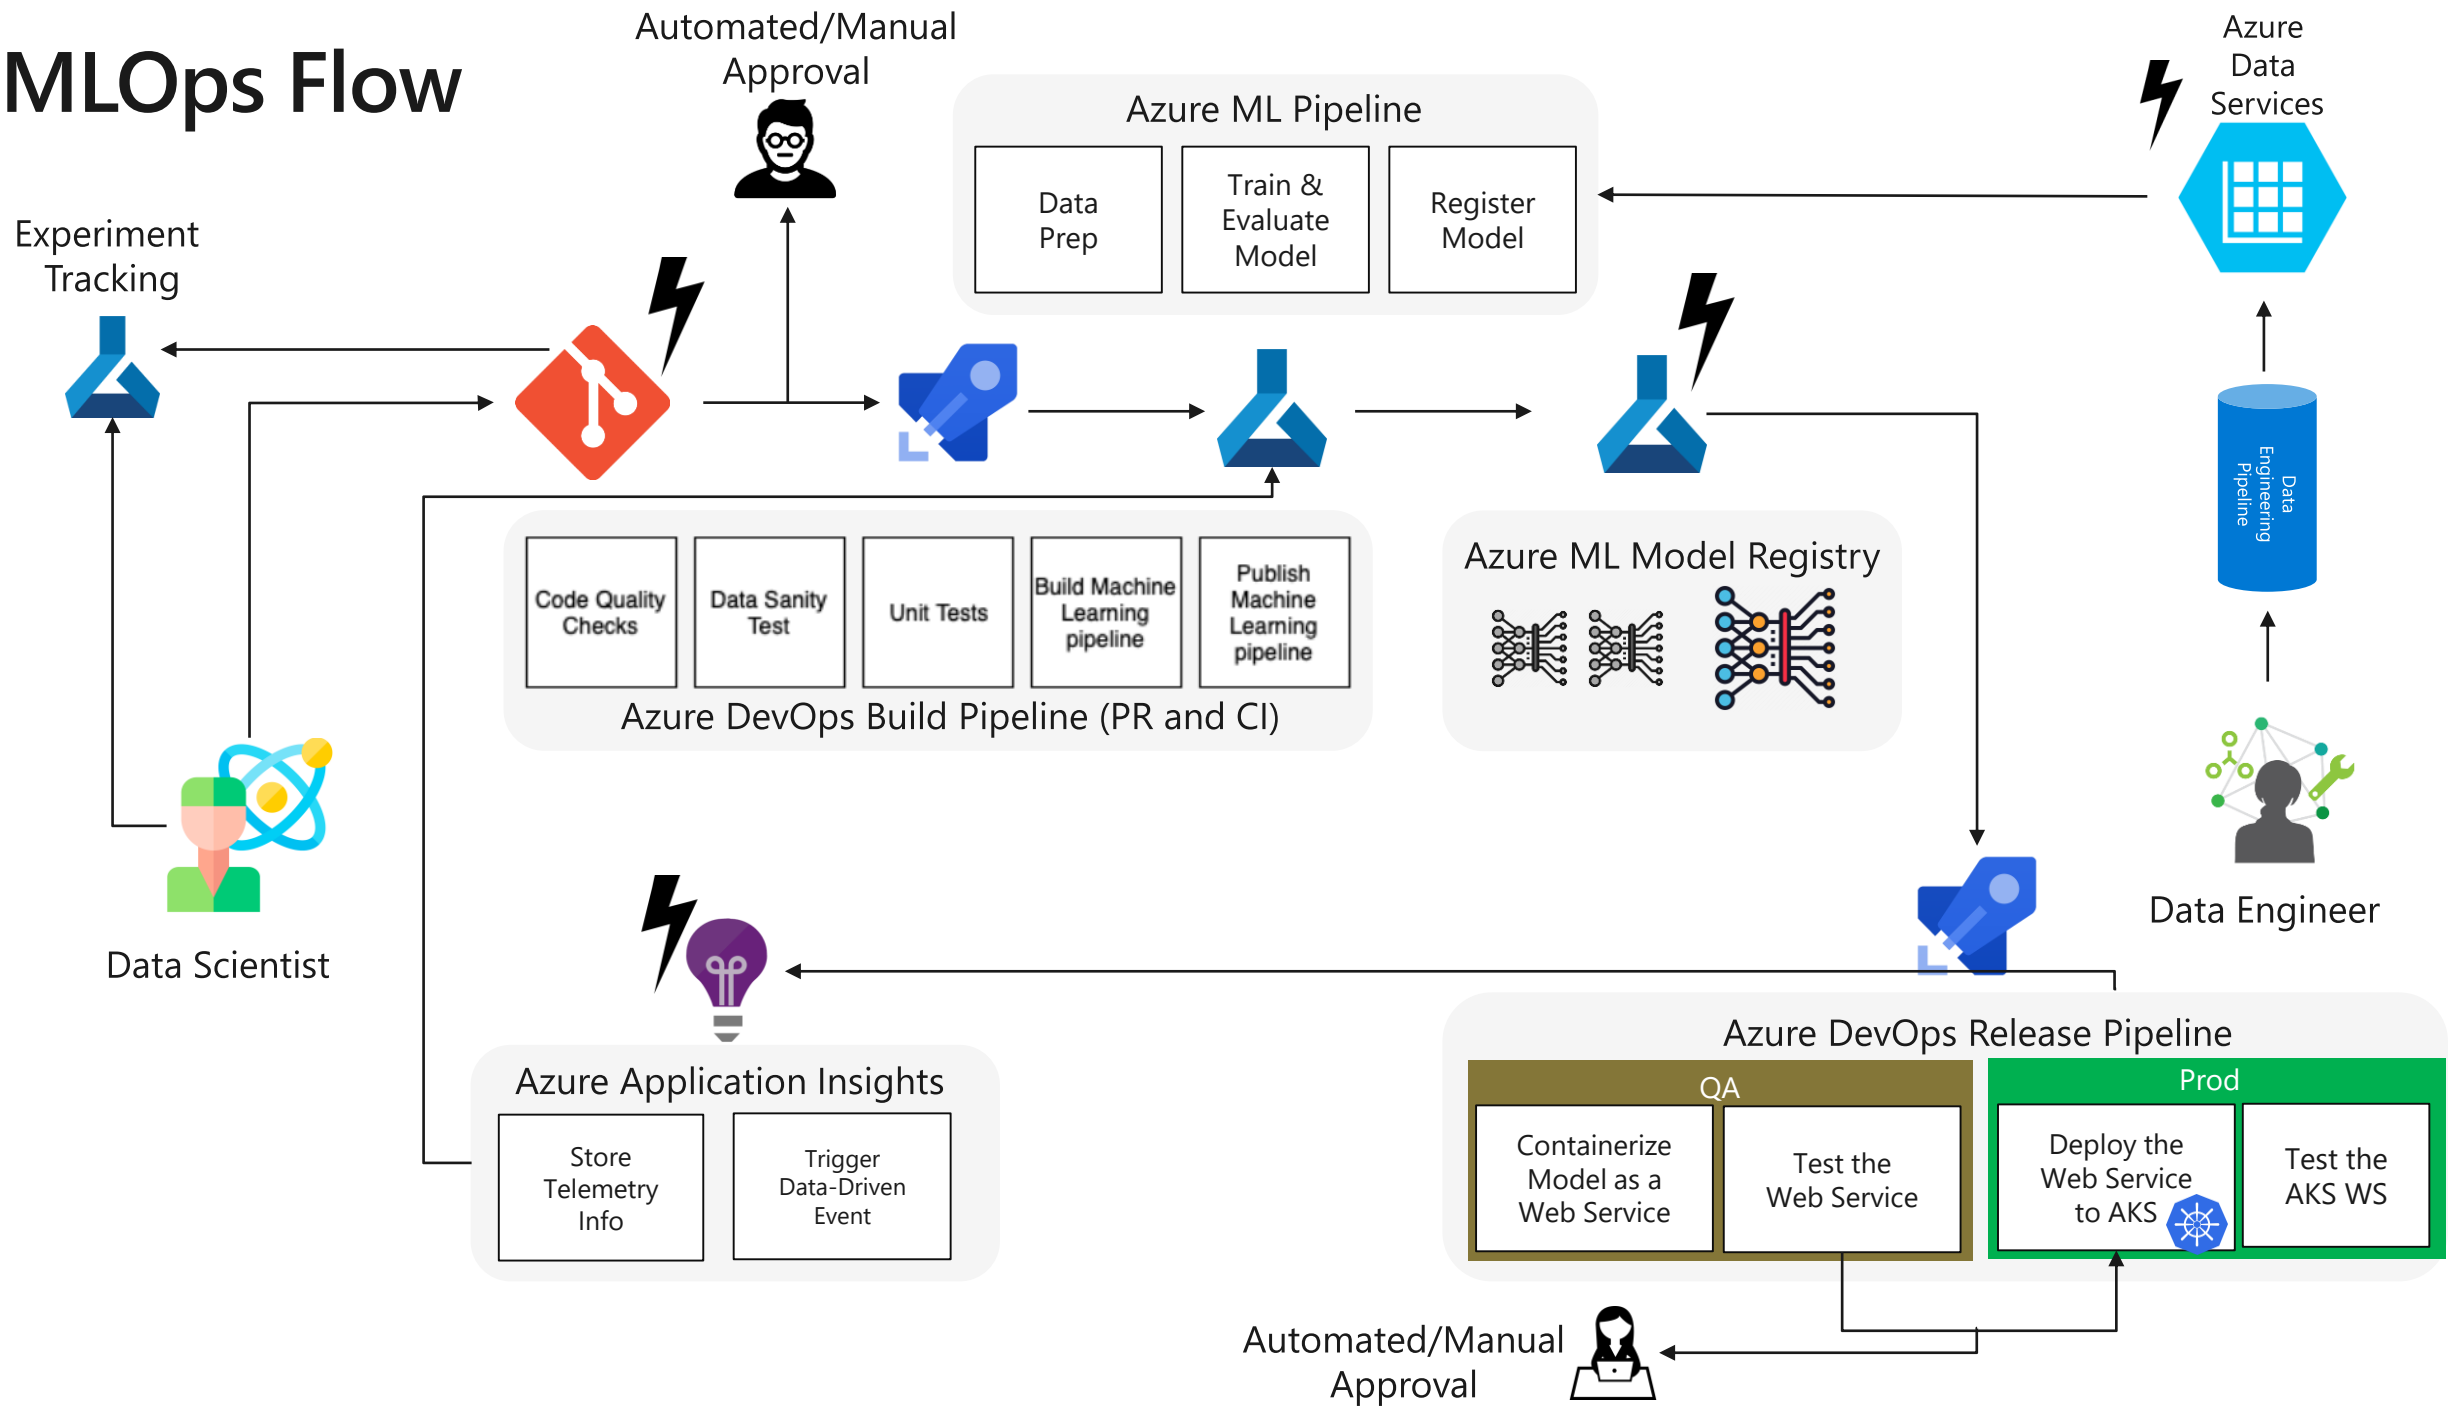


Model reuse is different than software reuse, as models must be tuned based on input data / scenario. To reuse a model you may need to fine-tune / transfer learn on it (meaning you need the training pipeline)



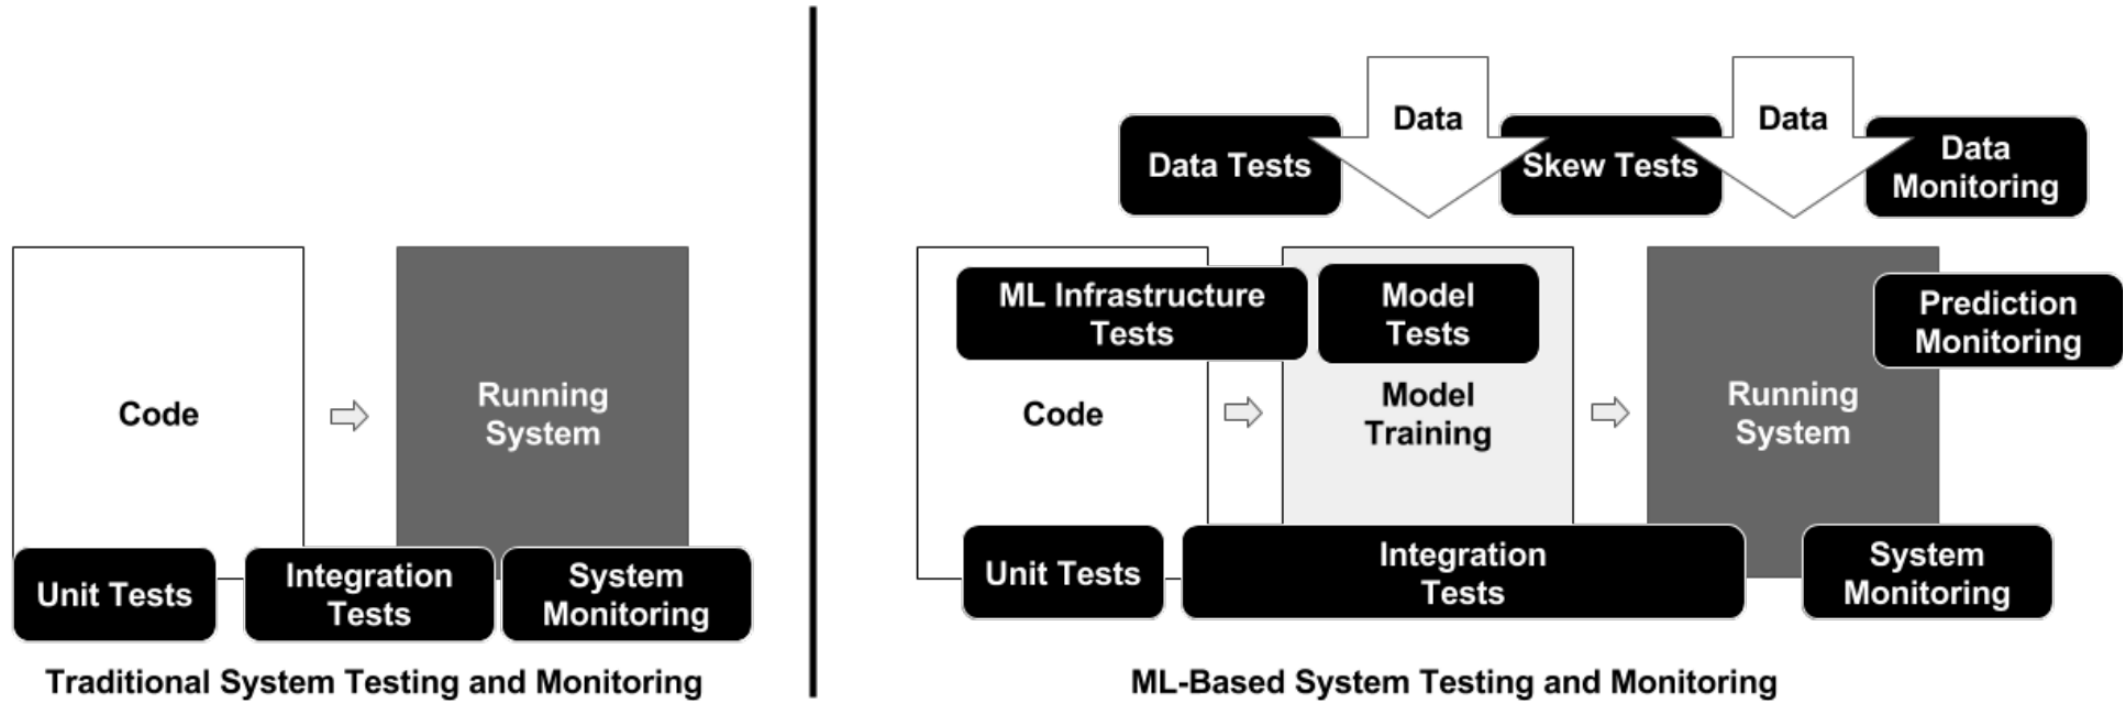
Model performance tends to decay over time & you need the ability to retrain them on demand to ensure they remain useful in a production context.

MLOps Flow



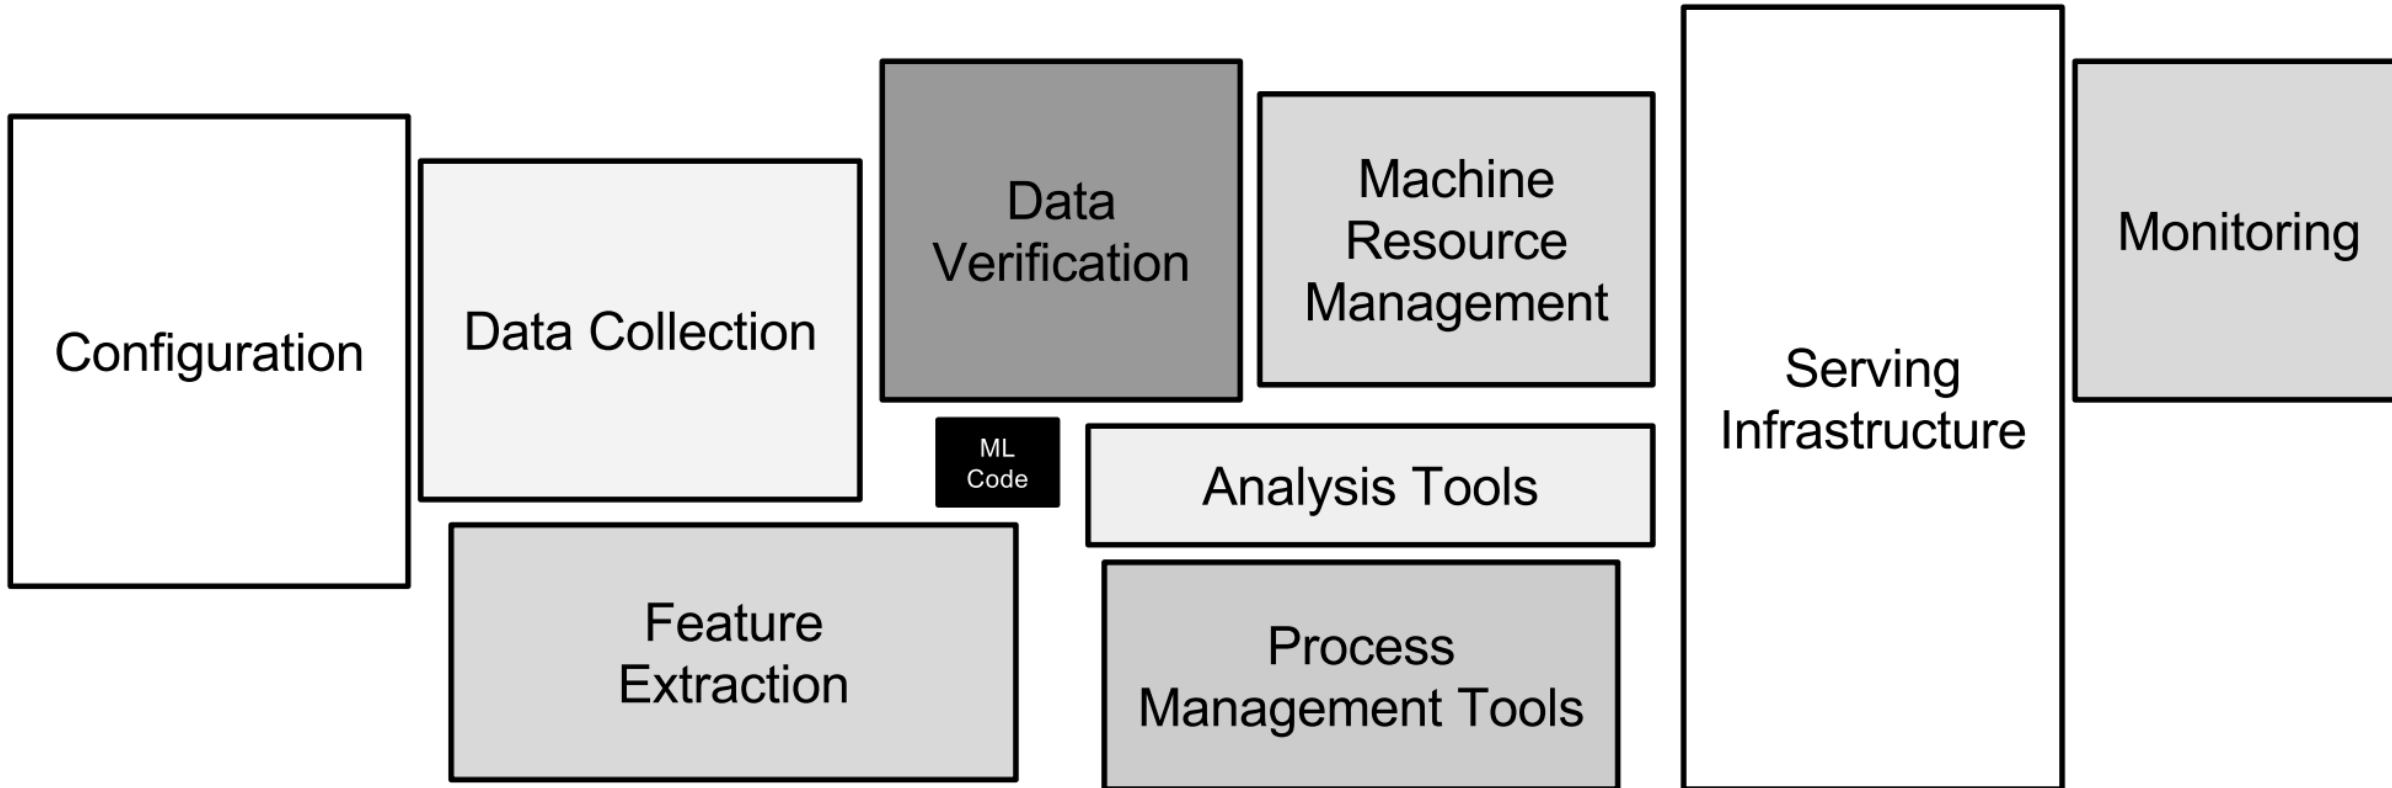
Why is MLOps hard?

Traditional vs. ML infused systems



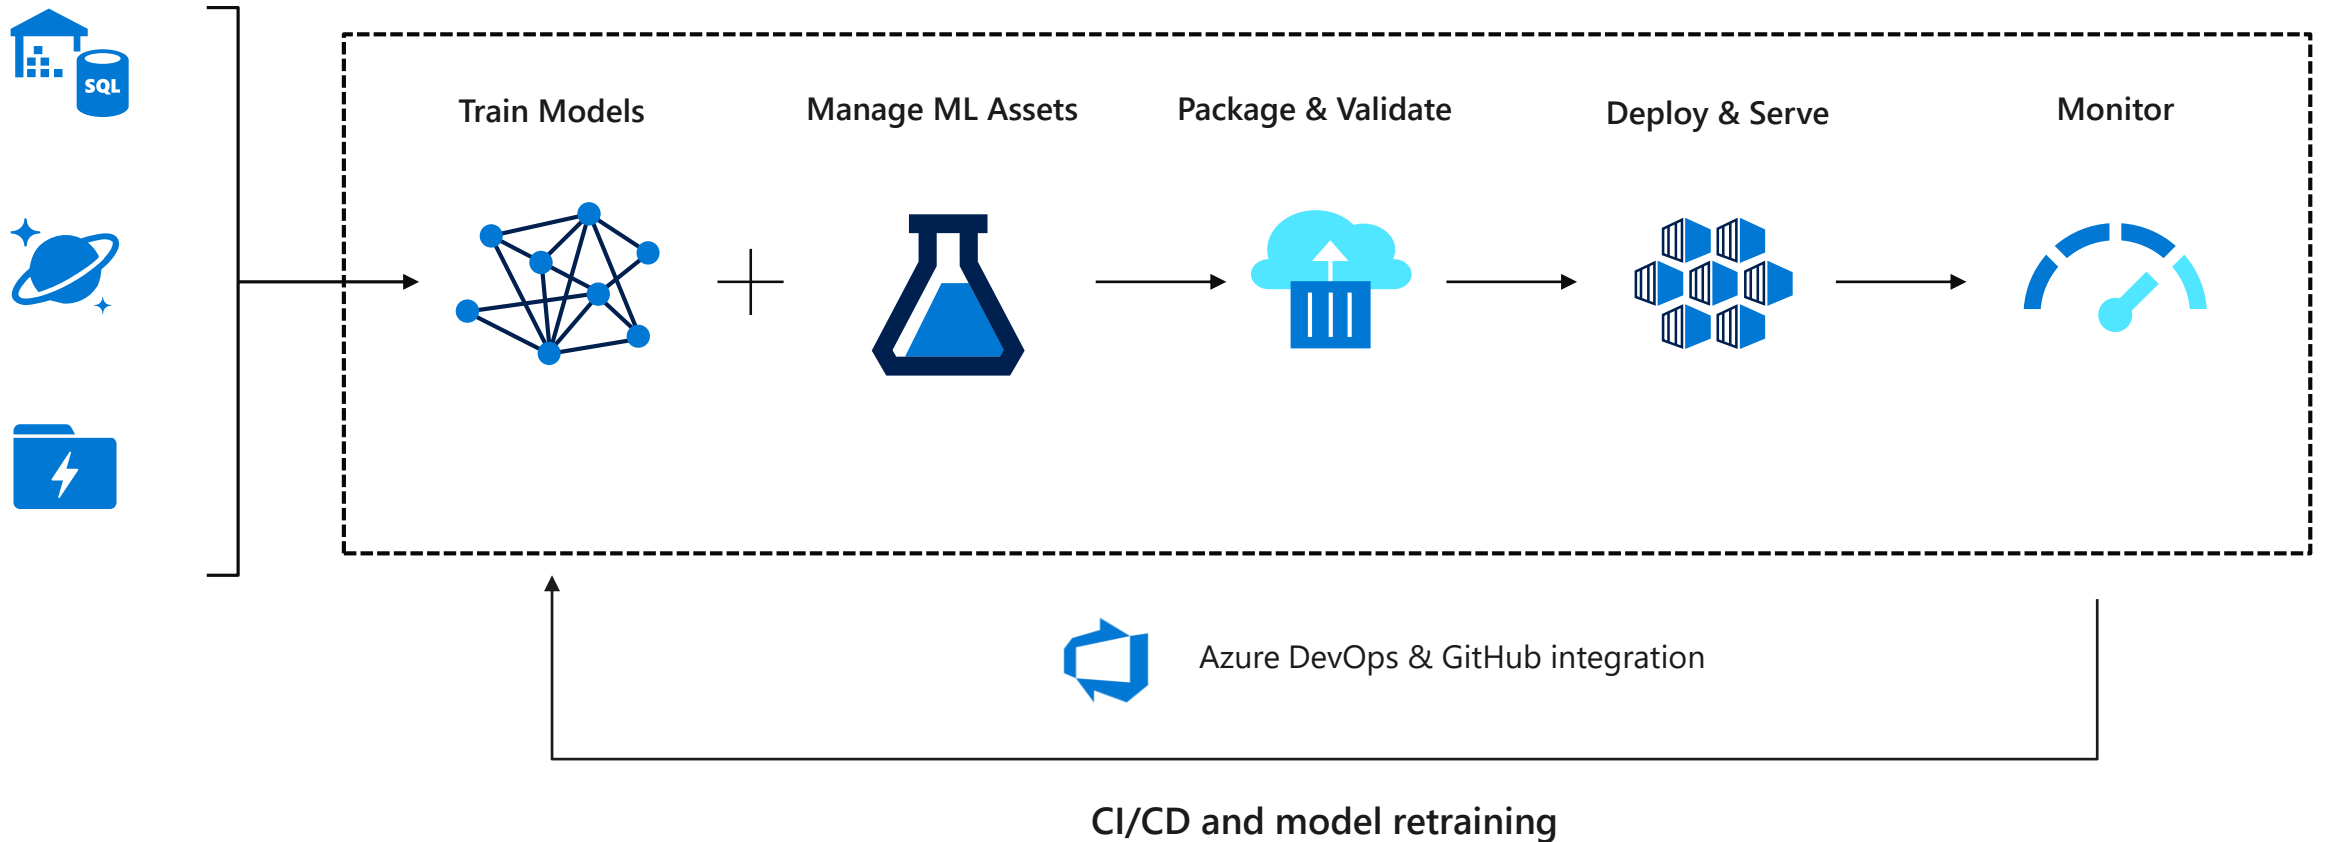
ML introduces two new assets into the software development lifecycle – **data** and **models**.

More assets & process to manage



Bring your ML workflows to production with MLOps

Azure Machine Learning service



What is MLOps on Azure?

What problems does MLOps solve?



Reproducibility of models and predictions



E2E operationalization & automation of ML lifecycle



Collaboration within / across teams on ML workflow



Enterprise readiness – security, governance, compliance, IaC

Azure ML Platform - MLOps Capabilities

Manage Assets



Infrastructure (Compute, Storage)



Code



Datasets



Environments



Runs



Models

Understand Model Behavior



Model Explainability



Model Profiling



Model Validation



Model Deployment

Enable Enterprise ML Lifecycle Management



Data Pipelines




ML Pipelines




Release Pipelines



Data-Driven Eventing



ML Lineage & Audit Trail



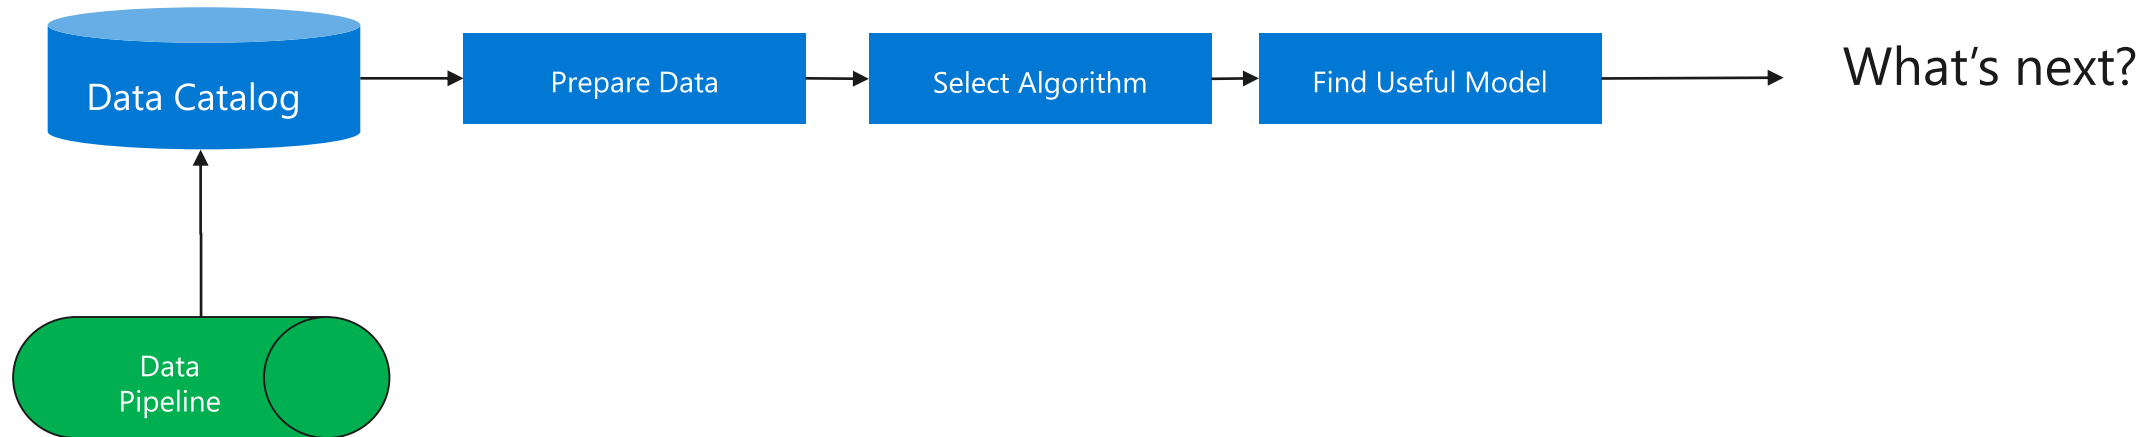
ML Workflow – Process Maturity Model

Phase 0 - Experimentation

Interactive, exploratory, get to something useful.



Data Scientist

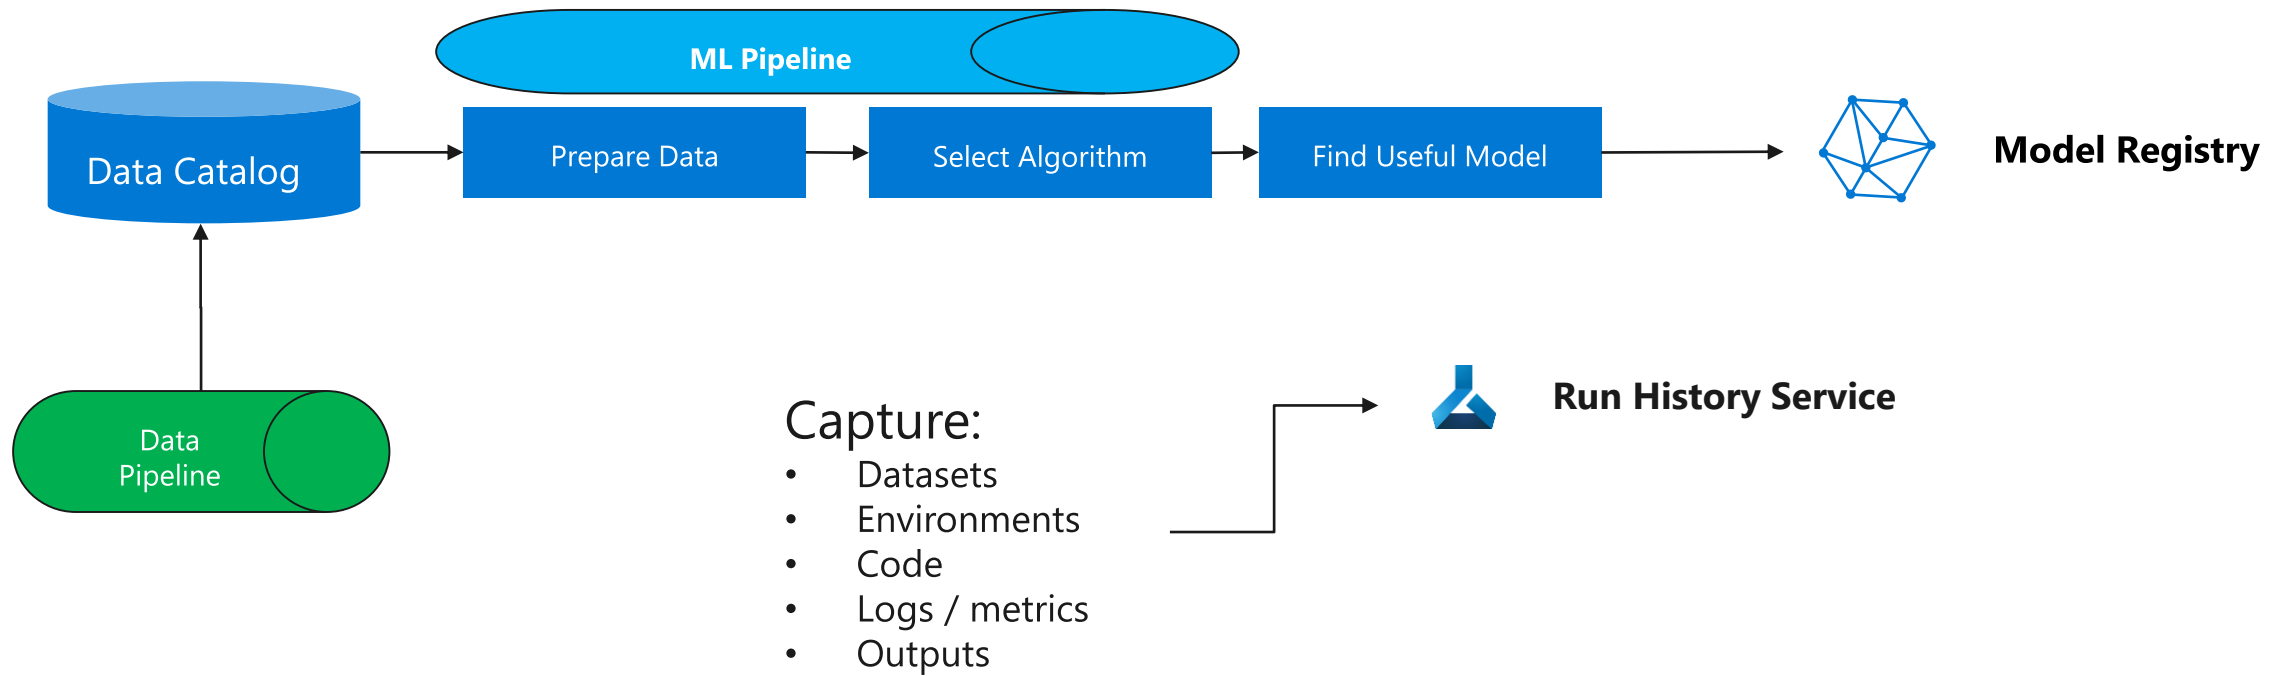


Phase 1A – Reproduce Model

Version code, data, ensure model can be recreated.

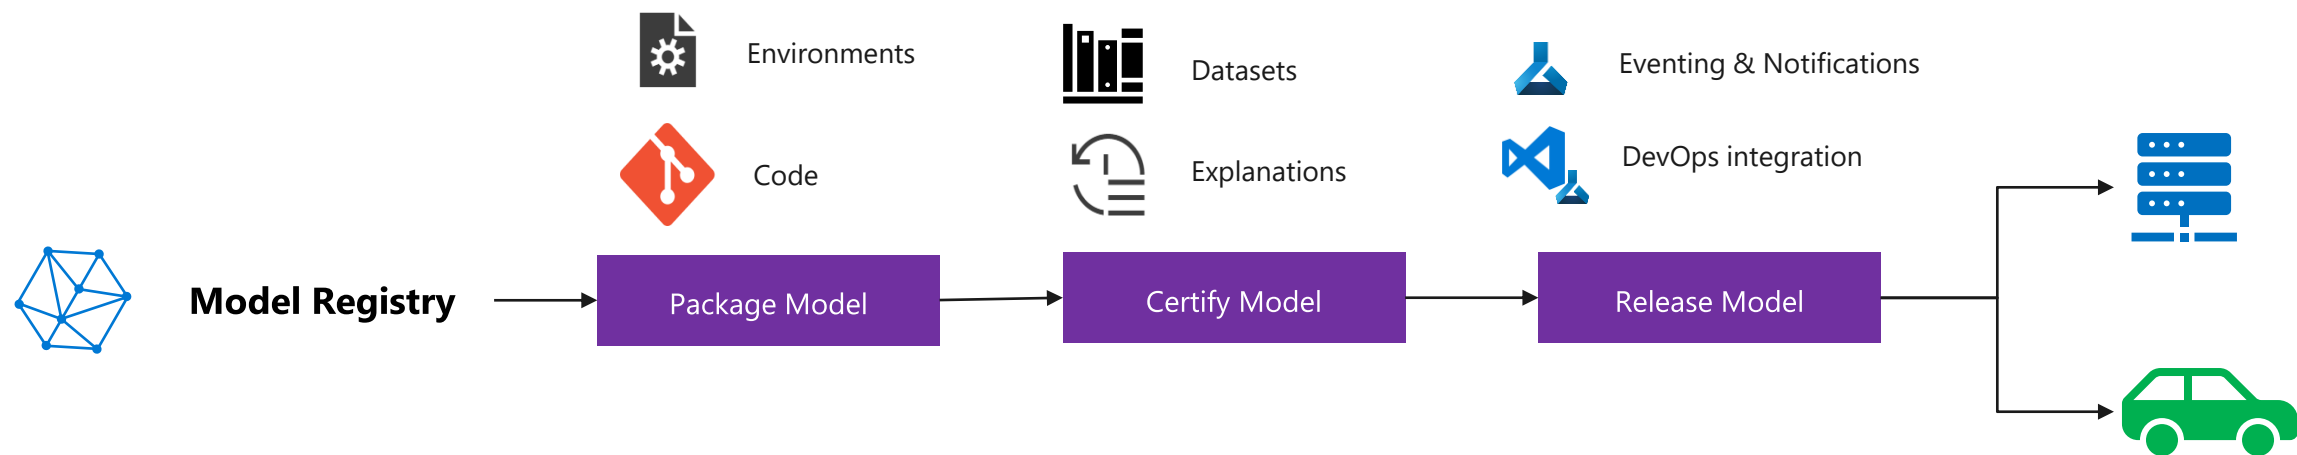


Data Scientist

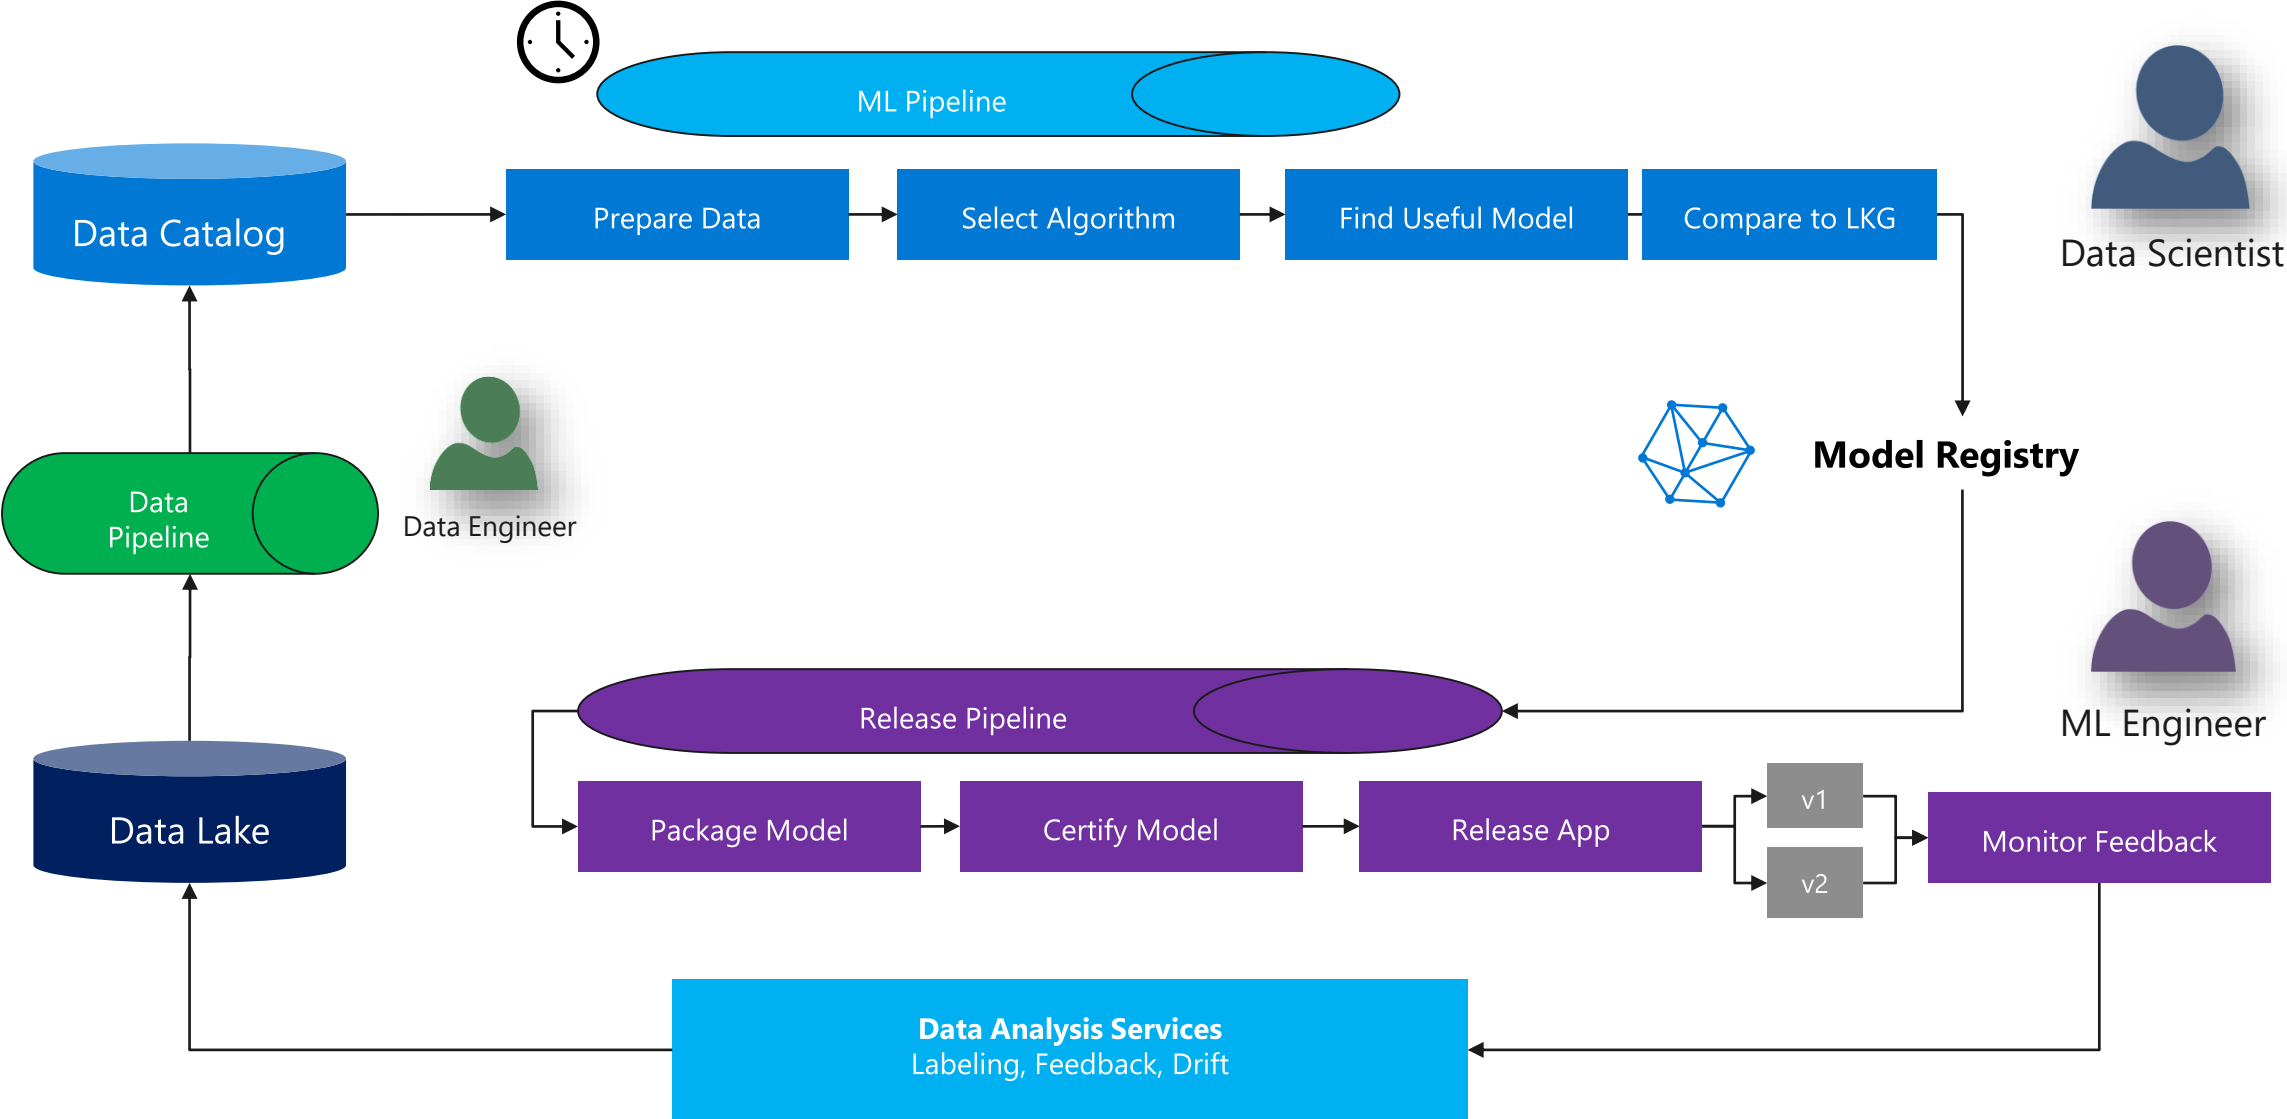


Phase 1B – Deploy Model

Package, certify, deploy



Phase 2 – Automate ML Lifecycle



What do we support?

Asset Management

- Code
- Datasets
- Experiments
- Models

Azure ML Services

- Model packaging
- Model profiling
- Model deployment

E2E Automation w/ CLI + DevOps integration

E2E Audit Trail through asset management

Asset Management

Microsoft Azure | Machine Learning

Home > Models

Models

Register Model

Archive

Refresh

Compare

...

Filter models...

Name == my-model

Type == Script Run

Add filter

Name	Version	Experiment	Run	Created by	Created Time	Tags
my-model-1	1	MyExperiment	Run 42	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-2	1	MyExperiment	Run 44	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-3	1	MyExperiment	Run 47	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-4	1	MyExperiment	Run 49	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-5	1	MyExperiment	Run 52	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-6	1	MyExperiment	Run 54	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-7	1	MyExperiment	Run 63	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-8	1	MyExperiment	Run 42	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-9	1	MyExperiment	Run 42	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression
my-model-10	1	MyExperiment	Run 42	Alan Woo	Jun 13, 2019 3:46 PM	area:diabetes type:regression

< Prev

Next >

Assets

Datasets?

Experiments

Pipelines

Models

Deployments

Modules

Management

Compute

Environments

Datastores

Data Labeling

Home

Authoring

Automated ML

Visual Interface

Notebooks and Files

someone@contoso.com
CONTOSO

Asset Management

1. Git Integration

Run Properties

Status	Completed
Start Time	5/5/2019 8:40:54 PM
Duration	0:08:56
Run Id	generate-attribution-explainer_1557114046_a98b1904
Arguments	N/A

Output Logs azureml-logs/80_driver_log.txt ☒ Auto-switch

100% 294/294 [01:02<00:00, 4.79it/s]

The experiment completed successfully. Finalizing run...
Logging experiment finalizing status in history service
Cleaning up all outstanding Run operations, waiting 300.0 seconds
5 items cleaning up...
Cleanup took 0.004108428955078125 seconds

Run is completed.

[Click here to see the run in Azure portal](#)

```
] : { 'runId': 'generate-attribution-explainer_1557114046_a98b1904',  
      'target': 'cpu',  
      'status': 'Completed',  
      'startTimeUtc': '2019-05-06T03:45:53.894802Z',  
      'endTimeUtc': '2019-05-06T03:49:50.583977Z',  
      'properties': { 'azureml.runsource': 'experiment',  
                     'AzureML.DerivedImageName': 'azureml/azureml_11a26c41b7e5439822ffd7fb19be3eaa',  
                     'ContentSnapshotId': 'e413abfe-0966-4960-b556-e1637ae1a53f',  
                     'azureml.git.repository_uri': 'https://github.com/Microsoft/MLOps_Hub.git',  
                     'azureml.git.branch': 'master',  
                     'azureml.git.commit': 'fe7d29bb1693047168d10c82949002aa9701037a',
```

Asset Management

2. Dataset Integration

In [82]: `run.get_details()`


```
'endTimeUtc': '2019-05-07T12:34:25.782485Z',  
'properties': {'azureml.runsource': 'experiment',  
'AzureML.DerivedImageName': 'azureml/azureml_385d8adbecd8ef667594f38ca071b409',  
'ContentSnapshotId': 'ea6c785d-ed57-409a-91f9-0c9f28e6b3f8',  
'azureml.git.repository_uri': 'https://github.com/Microsoft/MLOps_Hub.git',  
'azureml.git.branch': 'master',  
'azureml.git.commit': 'fe7d29bb1693047168d10c82949002aa9701037a',  
'azureml.git.dirty': 'True',  
'azureml.git.build_id': None,  
'azureml.git.build_uri': None,  
'mlflow.source.git.branch': 'master',  
'mlflow.source.git.commit': 'fe7d29bb1693047168d10c82949002aa9701037a',  
'mlflow.source.git.repoURL': 'https://github.com/Microsoft/MLOps_Hub.git',  
'azureml.dataset.get.ibm-attrition-data:1': '{"name": "ibm-attrition-data", "definition": "1", "snapshot": ""}',  
'model_type': 'classification',  
'explainer': 'tabular'},  
'runDefinition': {'script': 'train_explain.py',  
'arguments': [],  
'sourceDirectoryDataStore': None,  
'framework': 'Python',
```

Asset Management

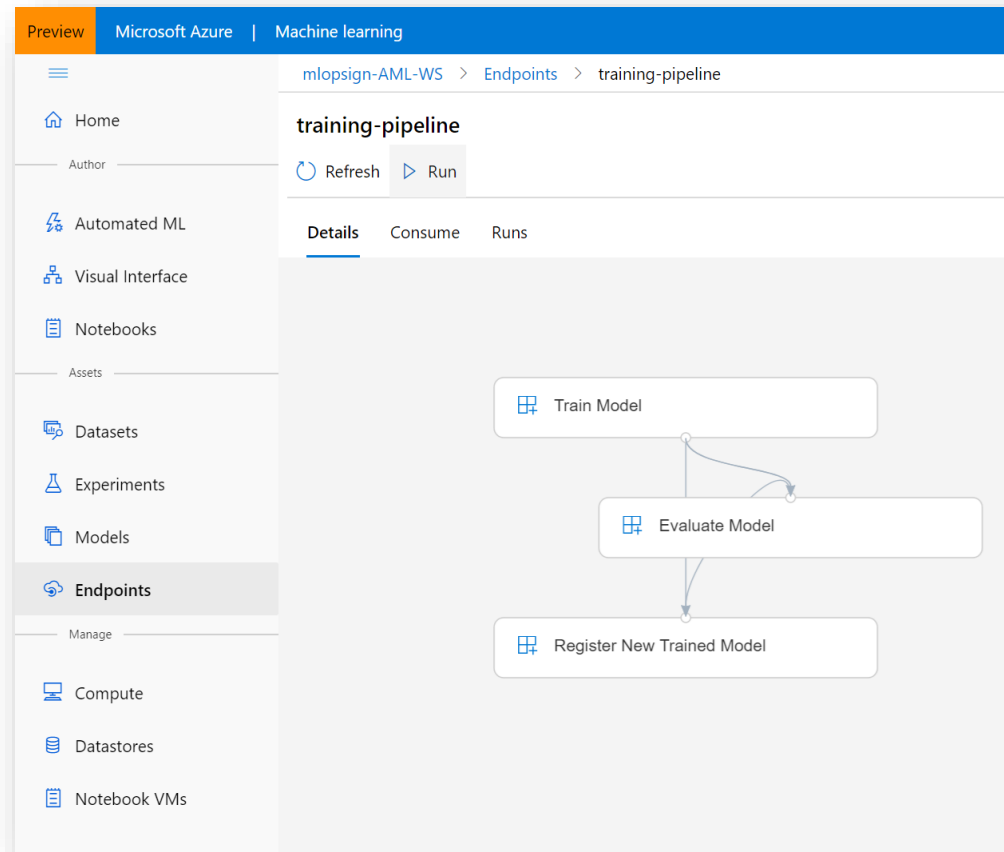
3. Environment

```
In [38]: run.get_environment()

Out[38]: {
  "name": "Experiment explainer-remote-run-on-amlcompute Environment",
  "version": "Autosave_2019-05-04T03:04:54Z_e5c501bb",
  "environmentVariables": {
    "EXAMPLE_ENV_VAR": "EXAMPLE_VALUE"
  },
  "python": {
    "userManagedDependencies": false,
    "interpreterPath": "python",
    "condaDependenciesFile": null,
    "baseCondaEnvironment": null,
    "condaDependencies": {
      "name": "project_environment",
      "dependencies": [
        "python=3.6.2",
        {
          "pip": [
            "--index-url https://azuremlsdktestpypi.azureedge.net/AzureML-Contrib-Explain-Model-Gated/3010237",
            "--extra-index-url https://pypi.python.org/simple",
            "sklearn_pandas",
            "azureml-defaults<0.1.50",
            "azureml-contrib-explain-model<0.1.50",
            "azureml-core<0.1.50",
            "azureml-telemetry<0.1.50",
            "azureml-explain-model<0.1.50"
          ]
        }
      ],
      "scikit-learn"
    }
  },
  "docker": {
    "enabled": true,
    "baseImage": "mcr.microsoft.com/azureml/base:intelmpi2018.3-ubuntu16.04",
    "sharedVolumes": true,
    "gpuSupport": false,
    "shmSize": "1g",
    "arguments": [],
    "baseImageRegistry": {
      "address": null,
      "username": null,
      "password": null
    }
  },
  "spark": {
    "repositories": [
      "https://mmlspark.azureedge.net/maven"
    ],
    "packages": [
      {
        "group": "com.microsoft.ml.spark",
        "artifact": "mmlspark_2.11",
        "version": "0.12"
      }
    ],
    "precachePackages": true
  },
}
```



Automated Training



```
from azureml.core.datastore import Datastore

datastore = Datastore(workspace=ws, name="workspaceblobstore")

schedule = Schedule.create(workspace=ws, name="My_Schedule",
                           pipeline_id=pub_pipeline_id,
                           experiment_name='Schedule_Run',
                           datastore=datastore,
                           wait_for_provisioning=True,
                           description="Schedule Run")
#polling_interval=5, use polling_interval to specify how often to poll for blob
additions/modifications. Default value is 5 minutes.
#path_on_datastore="file/path") use path_on_datastore to specify a specific fold
er to monitor for changes.
```


Automated Deployment

Pipeline Tasks Variables Retention Options History

Artifacts | + Add

Stages | + Add

_ci-build

_sklearn_regression_model.pkl

QA (ACI)
1 job, 1 task

Continuous deployment trigger

☒ Enabled

Enabling the trigger will create a new release every time a new AzureML model is registered

Pipeline Tasks Variables Retention Options History

QA (ACI)
Deployment process

Agent job
Run on agent

Azure ML Model Deploy

Display name *

Azure ML Model Deploy

☐ Validate and Profile ⓘ

Model Information ^

Azure ML Workspace * ⓘ | Manage ↗

mlops-AML-WS ⓘ

Scoped to resource 'workspaces' in resource group 'mlops-AML-RG'

Model Source * ⓘ

☒ Model Artifact ☐ Manual Specify

Inference config Path * ⓘ

\$(System.DefaultWorkingDirectory)/_ci-build/mlops-pipelines/code/scoring/inference_config.yml ⓘ

Deployment Information ^

Model Deployment Target * ⓘ

Azure Container Instance

Deployment Name * ⓘ

mlopspython-aci

Deployment configuration file * ⓘ

\$(System.DefaultWorkingDirectory)/_ci-build/mlops-pipelines/code/scoring/deployment_config_aci.yml ⓘ

☒ Overwrite existing deployment ⓘ

E2E Automation

Azure ML CLI

```
az ml model deploy -n mydetector -m videodetector:1 -  
-ic inferenceConfig.json --dc deploymentConfig.json
```

Triggers to automate model deployment using the Machine Learning extension

Continuous deployment trigger

☒ Enabled

① Enabling the trigger will create a new release every time a new AzureML model is registered



Machine Learning

Microsoft DevLabs | 630 installs | ★★★★★ (0) | Free

Submit experiments from a DevOps Pipeline, track code from Azure Repos or GitHub, trigger release pipelines when an ML model is registered, and automate ML deployments using Azure Pipelines.

Get it free

Integrated with Azure DevOps

A screenshot of the Azure DevOps Release Pipeline interface. The top navigation bar shows 'New release pipeline > Release-15'. Below this, there are tabs for 'Pipeline', 'Variables', and 'History'. The 'Pipeline' tab is active, showing a 'Release' section with a 'Continuous deployment' trigger for 'Microsoft.VisualStudio' on '5/2/2019, 6:26 PM'. The 'Artifacts' section lists two artifacts: '_Build-MLOps-Demo' (dcee1f07) and '_mymodel' (4). The 'Stages' section shows 'Stage 1' with a 'Succeeded' status and a timestamp of 'on 5/2/2019, 6:34 PM'.

Demo video of MLOps E2E

- <https://www.youtube.com/watch?v=1XJwcfQP7YI>

Key Takeaways

Better together: ML + DevOps mindset

MLOps provides structure for building, deploying and managing and an enterprise-ready AI application lifecycle

MLOps enhances delivery

Adoption will increase the agility, quality and delivery of AI project teams.

More than technology

MLOps is a conversation about people, process and technology
AI principles and practices need to be understood by all roles

